

Joining and pivoting tables

Raphael Eisenhofer

2022_06_02

Outline for today:

- 1. Introduction to tidy data
- 2. How to pivot tables longer or wider
- 3. Joining tables

I. Introduction to tidy data



Tidy data

- Hadley Wickham: see original paper here: <https://vita.had.co.nz/papers/tidy-data.pdf>
- Tidy data is a consistent way of representing data
 - Benefits include:
 - Easier to interpret
 - Easier to learn tools that manipulate it (e.g. dplyr)
 - Faster (takes advantage of R's vectorized nature)
- The three rules of tidy data:
 - **1: Each variable forms a column**
 - **2: Each observation forms a row**
 - **3: Each value must have its own cell**



Tidy data example:

- <https://about.dataclassroom.com/blog/keep-your-data-tidy>

Not tidy!

Day	Plant A height (cm)	Plant B height (cm)
1	0.7	1.5
2	1.0	0.7
3	1.5	0.9
4	1.8	1.3
5	2.2	1.8



Tidy

Day	Plant	Height (cm)
1	A	0.7
2	A	1.0
3	A	1.5
4	A	1.8
5	A	2.2
1	B	1.5
2	B	0.7
3	B	0.9
4	B	1.3
5	B	1.8



Tidy data example:

- <https://r4ds.had.co.nz/tidy-data.html>

Tidy

country	year	key	value
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898
China	1999	cases	212258
China	1999	population	1272915272
China	2000	cases	213766
China	2000	population	1280428583

table2



Tidy data example:

- <https://r4ds.had.co.nz/tidy-data.html>

Tidy

country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

country	1999	2000
Afghanistan	745	2666
Brazil	37737	80488
China	212258	213766

table4

2. How to get tidy

Pivoting longer or wider

`pivot_longer()`



© @ketama.de.galgo.SPLITPICS UK

`pivot_wider()`





pivot_longer()

- **Syntax:**

- **x** = *dataframe*
- **cols** = *columns to make longer* [e.g. `c(1999, 2000)`]
 - can also use `!` (for not), e.g. `cols = !country`
- **names_to** = *name of new column*, e.g. “year”
 - Can choose multiple, but need to provide **names_sep** or **names_pattern**
- **values_to** = *name of new column to store values*, e.g. “cases”

Optional:

values_drop_na = TRUE

- Will remove redundant rows 😊

names_pattern = “pattern”

- Determines how name will be broken, can use regular expressions e.g. `(.*)_suffix`

country	year	cases	country	1999	2000
Afghanistan	1999	745	Afghanistan	745	2666
Afghanistan	2000	2666	Brazil	37737	80488
Brazil	1999	37737	China	212258	213766
Brazil	2000	80488			
China	1999	212258			
China	2000	213766			

table4

pivot_wider()



- **Syntax:**

- **x** = *dataframe*
- **names_from** = *name of column to widen, e.g. “key”*
 - Can choose multiple, but need to provide **names_sep** or **names_pattern**
- **values_from** = *name of column that has values e.g. “value”*

- **Optional:**

- **values_fill** = value for missing values e.g. 0

country	year	key	value
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898
China	1999	cases	212258
China	1999	population	1272915272
China	2000	cases	213766
China	2000	population	1280428583

table2



separate()

- **Syntax:**

- **col** = *name of the column* (e.g. rate)
- **into** = *name of new columns* – e.g. c(“cases”, “population”)
- **sep** = *separator character* (e.g. “/”) – by default it will guess non-alphanumeric chars
- **convert** = TRUE or FALSE (default behaviour is to leave column types as is – ‘chr’, so by using TRUE it will guess and change to an integer for us!)

country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

table3

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

unite()



- **Syntax:**

- **col** = *name of new column* e.g. year
- **...** = *columns you want to join* e.g. century, year
- **sep** = *delimiter to put between columns* (use “” for no separator)

country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

country	century	year	rate
Afghanistan	19	99	745 / 19987071
Afghanistan	20	0	2666 / 20595360
Brazil	19	99	37737 / 172006362
Brazil	20	0	80488 / 174504898
China	19	99	212258 / 1272915272
China	20	0	213766 / 1280428583

table6

3. Joining tables



The problem

	A	B	C
1	genome	completeness	contamination
2	bin_m1.cct123	99.23	0.35
3	bin_m1.mtb106	93.55	0.05
4	bin_m1.mtb147	96.29	4.36
5	bin_m1.mtb2	89.97	0.88
6	bin_m1.mxb107_sub	36.3	15.58
7	bin_m1.vmb35	96.43	0.95
8	bin_m1.vmb46	95.49	1.69
9	bin_m10.cct0	87.8	0.81
10	bin_m10.cct151	89.81	2.75
11	bin_m10.cct29	61.97	1.59
12	bin_m10.cct54_sub	91.21	0
13	bin_m10.cct7_sub	85.71	1.16
14	bin_m10.mtb109	97.55	1.13
15	bin_m10.mtb12	91.78	3.02



	A	B	C	D	E	F	G	H
1	user_genome	classification	fastani_refer	fastani_refer	fastani_taxo	fastani_ani	fastani_af	closest_place
2	bin_m1.cct123	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	N/A
3	bin_m1.mtb106	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	GCA_003979
4	bin_m1.mtb147	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	GCF_011958
5	bin_m1.mtb2	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	N/A
6	bin_m1.mxb107_sub	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	N/A
7	bin_m1.vmb35	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	GCA_002298
8	bin_m1.vmb46	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	GCF_011959
9	bin_m10.cct0	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	N/A
10	bin_m10.cct151	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	N/A
11	bin_m10.cct29	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	N/A
12	bin_m10.cct54_sub	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	N/A
13	bin_m10.cct7_sub	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	GCA_013316
14	bin_m10.mtb109	d__Bacteria;	GCF_003762	95	d__Bacteria;	96.83	0.89	GCF_003762
15	bin_m10.mtb12	d__Bacteria;	N/A	N/A	N/A	N/A	N/A	GCA_900554

- What if they are not ordered in the same way?
- What if tables do not have the same number of rows?



The simplest, but non-ideal way

- Simplest way: `bind_cols()`
- **N.B.** does not match tables by ID, so you have to ensure that both tables are sorted correctly!

X				y				
A	B	C		E	F	G		
a	t	1		a	t	3		
b	u	2		b	u	2		
c	v	3		d	w	1		

+

A	B	C	E	F	G
a	t	1	a	t	3
b	u	2	b	u	2
c	v	3	d	w	1

=

A	B	C	E	F	G
a	t	1	a	t	3
b	u	2	b	u	2
c	v	3	d	w	1

Relational joining of tables



- We can do this a bit smarter by matching values in the rows between two tables (relational)
- E.g. say that both tables have the same grouping variable:
 - sample_id
 - genome_name
 - etc.

A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA

left_join()

A	B	C	D
a	t	1	3
b	u	2	2
d	w	NA	1

right_join()

A	B	C	D
a	t	1	3
b	u	2	2

inner_join()

A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA
d	w	NA	1

full_join()



Relational joining of tables

Example: `left_join(x = dataframe1,`
 `y = dataframe2,`
 `by = "col name" OR c("col name1", "colname2"),`
 `)`

- If both tables have the same column name for grouping, `by = "column name"`
- If both tables have different names (e.g. `x = grouping_var_x` & `y = grouping_var_y`), then use a vector: `by = c("grouping_var_x", "grouping_var_y")`
- If you have multiple tables, you can pipe `%>%` multiple `*_join()` together!

Exercise time!

- See the exercises in “Exercises.R”



Summary

- Tidy data is consistent, easier to learn to manipulate, and faster
- Dplyr has some really nice tools for manipulating tables to get what you want
 - `pivot_longer()` and `pivot_wider()`
 - `spread()` and `gather()`
- Joining tables in a relational fashion saves time, and is less prone to error!

